
Zhe Gan

Microsoft Advanta-B, 6227,
4200 150th Ave NE,
Redmond, WA 98052

Phone: Provided upon request
Email: zhe.gan@microsoft.com
Homepage: <http://zhegan27.github.io/>

Research Interests

I am a Senior Researcher at Microsoft Dynamics 365 AI Research, primarily working on Vision-and-Language Representation Learning, Generative Pre-training, and Adversarial Machine Learning. I also have broad interests on other machine learning topics.

Education

- Duke University, Durham, NC
Ph.D., Electrical and Computer Engineering 09/2013 - 03/2018
- Peking University, Beijing, China
M.S., Electrical Engineering 09/2010 - 07/2013
B.S., Electrical Engineering 09/2006 - 07/2010

Experience

- **Microsoft Cloud and AI** 04/2018 - present
Senior Researcher. Manager: Jingjing Liu, MS D365 AI
Vision-and-Language Representation Learning, Generative Pre-training, Self-supervised Learning, and Adversarial Machine Learning
- **Information Initiative at Duke (iiD)** 09/2013 - 03/2018
Research Assistant. Advisor: Prof. Lawrence Carin
(i) Deep Bayesian Learning: developing deep generative models for computer vision and natural language processing applications, including VAE and GAN
(ii) Bayesian Deep Learning: designing stochastic gradient variational inference algorithms and stochastic gradient MCMC methods for scalable Bayesian inference
- **Microsoft Research Redmond** 05/2017 - 08/2017
Research Intern. Advisor: Xiaodong He, Lihong Li, Ph.D
Deep reinforcement learning for vision and language intelligence, with focus on the visual storytelling task.
- **Microsoft Research Redmond** 05/2016 - 08/2016
Research Intern. Advisor: Xiaodong He, Jianfeng Gao, Li Deng, Ph.D
(i) image captioning: using deep learning techniques to improve the state-of-the-art of image and video captioning.
(ii) deep conflation: using deep learning techniques to implement conflation for business data analytics.
- **Adobe Research** 06/2015 - 09/2015
Data Scientist Intern. Advisor: Hung Bui, Ph.D
Recurrent neural networks (RNN) for NLP applications, including sentence classification, sentence retrieval and sentence generation

Publications

arXiv preprints

1. **Z. Gan**, Y.-C. Chen, L. Li, C. Zhu, Y. Cheng and J. Liu "Large-Scale Adversarial Training for Vision-and-Language Representation Learning", *arXiv preprint arXiv:2006.06195*
2. C. Zhu, Y. Cheng, **Z. Gan**, F. Huang, J. Liu and T. Goldstein "Adaptive Learning Rates with Maximum Variation Averaging", *arXiv preprint arXiv:2006.11918*
3. L. Li*, Y.-C. Chen*, Y. Cheng, **Z. Gan**, L. Yu and J. Liu "HERO: Hierarchical Encoder for Video+Language Omni-representation Pre-training", *arXiv preprint arXiv:2005.00200*
4. Y. Zhang*, G. Wang*, C. Li, **Z. Gan**, C. Brockett and B. Dolan "POINTER: Constrained Text Generation via Insertion-based Generative Pre-training", *arXiv preprint arXiv:2005.00558*
5. Y. Cheng, **Z. Gan**, Y. Zhang, O. Elachqar, D. Li and J. Liu "Contextual Text Style Transfer", *arXiv preprint arXiv:2005.00136*
6. S. Dai, **Z. Gan**, Y. Cheng, C. Tao, L. Carin and J. Liu "APo-VAE: Text Generation in Hyperbolic Space", *arXiv preprint arXiv:2005.00054*
7. Y. Fang, S. Sun, **Z. Gan**, R. Pillai, S. Wang and J. Liu "Hierarchical Graph Network for Multi-hop Question Answering", *arXiv preprint arXiv:1911.03631*
8. S. Dai, Y. Cheng, Y. Zhang, **Z. Gan**, J. Liu and L. Carin "Contrastively Smoothed Class Alignment for Unsupervised Domain Adaptation", *arXiv preprint arXiv:1909.05288*

2021

1. W. Chen, **Z. Gan**, L. Li, Y. Cheng, W. Wang and J. Liu "Meta Module Network for Compositional Visual Reasoning", *Winter Conf. on Applications of Computer Vision (WACV)*, 2021 **Oral**

2020

1. J. Cao, **Z. Gan**, Y. Cheng, L. Yu, Y.-C. Chen and J. Liu "Behind the Scene: Revealing the Secrets of Pre-trained Vision-and-Language Models", *European Conf. on Computer Vision (ECCV)*, 2020 **Spotlight**
2. Y.-C. Chen*, L. Li*, L. Yu*, A. Kholy, F. Ahmed, **Z. Gan**, Y. Cheng and J. Liu "UNITER: UNiversal Image-TEXT Representation Learning", *European Conf. on Computer Vision (ECCV)*, 2020
3. Y. Cheng, **Z. Gan**, Y. Li, J. Liu and J. Gao "Sequential Attention GAN for Interactive Image Editing via Dialogue", *ACM International Conference on Multimedia (ACMMM)*, 2020
4. P. Cheng, W. Hao, S. Dai, J. Liu, **Z. Gan** and L. Carin "CLUB: A Contrastive Log-ratio Upper Bound of Mutual Information", *Int. Conf. Machine Learning (ICML)*, 2020
5. L. Chen, **Z. Gan**, Y. Cheng, L. Li, L. Carin and J. Liu "Graph Optimal Transport for Cross-Domain Alignment", *Int. Conf. Machine Learning (ICML)*, 2020
6. J. Xu, **Z. Gan**, Y. Cheng and J. Liu "Discourse-Aware Neural Extractive Text Summarization", *Association for Computational Linguistics (ACL)*, 2020
7. Y. Chen, **Z. Gan**, Y. Cheng, J. Liu and J. Liu "Distilling Knowledge Learned in BERT for Text Generation", *Association for Computational Linguistics (ACL)*, 2020
8. R. Zhang, C. Chen, **Z. Gan**, W. Wang, D. Shen, G. Wang, Z. Wen and L. Carin "Improving Adversarial Text Generation by Modeling the Distant Future", *Association for Computational Linguistics (ACL)*, 2020
9. Y. Li, Y. Cheng, **Z. Gan**, L. Yu, L. Wang, J. Liu "BachGAN: High-Resolution Image Synthesis from Salient Object Layout", *Computer Vision and Pattern Recognition (CVPR)*, 2020
10. J. Liu, W. Chen, Y. Cheng, **Z. Gan**, L. Yu, Y. Yang, J. Liu "VIOLIN: A Large-Scale Dataset for Video-and-Language Inference", *Computer Vision and Pattern Recognition (CVPR)*, 2020
11. R. Zhang, C. Chen, **Z. Gan**, Z. Wen, W. Wang, L. Carin "Nested-Wasserstein Self-Imitation Learning for Sequence Generation", *Artificial Intelligence and Statistics (AISTATS)*, 2020
12. C. Zhu, Y. Cheng, **Z. Gan**, S. Sun, T. Goldstein and J. Liu "FreeLB: Enhanced Adversarial Training for Natural Language Understanding", *Int. Conf. Learning Representations (ICLR)*, 2020 **Spotlight**

13. W. Wang, H. Xu, **Z. Gan**, B. Li, G. Wang, L. Chen, Q. Yang, W. Wang and L. Carin “Graph-Driven Generative Models for Heterogeneous Multi-Task Learning”, *Proc. American Association of Artificial Intelligence (AAAI)*, 2020 **Spotlight**
14. J. Hu, Y. Cheng, **Z. Gan**, J. Liu, J. Gao and G. Neubig “What Makes A Good Story? Designing Composite Rewards for Visual Storytelling”, *Proc. American Association of Artificial Intelligence (AAAI)*, 2020 **Spotlight**

2019

1. W. Wang, C. Tao, **Z. Gan**, G. Wang, L. Chen, X. Zhang, R. Zhang, Q. Yang, R. Henao and L. Carin “Improving Textual Network Learning with Variational Homophilic Embeddings”, *Neural Information Processing Systems (NeurIPS)*, 2019
2. R. Zhang, C. Chen, **Z. Gan**, Z. Wen, W. Wang, L. Carin “Nested-Wasserstein Distance for Sequence Generation”, *Workshop on Bayesian Deep Learning, NeurIPS 2019*
3. S. Sun, Y. Cheng, **Z. Gan**, and J. Liu “Patient Knowledge Distillation for BERT Model Compression”, *Conf. on Empirical Methods in Natural Language Processing (EMNLP)*, 2019
4. H. Wang, **Z. Gan**, X. Liu, J. Liu, J. Gao and H. Wang “Adversarial Domain Adaptation for Machine Reading Comprehension”, *Conf. on Empirical Methods in Natural Language Processing (EMNLP)*, 2019
5. D. Li, Y. Zhang, **Z. Gan**, Y. Cheng, C. Brockett, M. Sun and B. Dolan “Domain Adaptive Text Style Transfer”, *Conf. on Empirical Methods in Natural Language Processing (EMNLP)*, 2019
6. M. Jiang, Q. Huang, L. Zhang, X. Wang, P. Zhang, **Z. Gan**, J. Diesner and J. Gao “TIGER: Text-to-Image Grounding for Image Caption Evaluation”, *Conf. on Empirical Methods in Natural Language Processing (EMNLP)*, 2019
7. L. Li, **Z. Gan**, Y. Cheng and J. Liu “Relation-Aware Graph Attention Network for Visual Question Answering”, *Int. Conf. on Computer Vision (ICCV)*, 2019
8. **Z. Gan**, Y. Cheng, A. Kholy, L. Li, J. Liu and J. Gao “Multi-step Reasoning via Recurrent Dual Attention for Visual Dialog”, *Association for Computational Linguistics (ACL)*, 2019
9. L. Ke, X. Li, Y. Bisk, A. Holtzman, **Z. Gan**, J. Liu, J. Gao, Y. Choi, and S. Srinivasa “Tactical Rewind: Self-Correction via Backtracking in Vision-and-Language Navigation”, *Computer Vision and Pattern Recognition (CVPR)*, 2019 **Oral**
10. Y. Li, **Z. Gan**, Y. Shen, J. Liu, Y. Cheng, Y. Wu, L. Carin, D. Carlson and J. Gao “StoryGAN: A Sequential Conditional GAN for Story Visualization”, *Computer Vision and Pattern Recognition (CVPR)*, 2019
11. W. Wang, **Z. Gan**, H. Xu, R. Zhang, G. Wang, D. Shen, C. Chen and L. Carin “Topic-Guided Variational Autoencoders for Text Generation”, *North American Chapter of the Association for Computational Linguistics (NAACL)*, 2019 **Oral**
12. L. Chen, Y. Zhang, R. Zhang, C. Tao, **Z. Gan**, H. Zhang, B. Li, D. Shen, C. Chen and L. Carin “Improving Sequence-to-Sequence Learning via Optimal Transport”, *Int. Conf. Learning Representations (ICLR)*, 2019
13. Q. Huang*, **Z. Gan***, A. Celikyilmaz, D. Wu, J. Wang and X. He “Hierarchically Structured Reinforcement Learning for Topically Coherent Visual Story Generation”, *Proc. American Association of Artificial Intelligence (AAAI)*, 2019 **Spotlight**

2018

1. Y. Zhang, M. Galley, J. Gao, **Z. Gan**, X. Li, C. Brockett and B. Dolan “Generating Informative and Diverse Conversational Responses via Adversarial Information Maximization”, *Neural Information Processing Systems (NeurIPS)*, 2018
2. L. Chen, S. Dai, C. Tao, D. Shen, **Z. Gan**, H. Zhang, Y. Zhang and L. Carin “Adversarial Text Generation via Feature-Mover’s Distance”, *Neural Information Processing Systems (NeurIPS)*, 2018
3. X. Zhang, R. Henao, **Z. Gan**, Y. Li and L. Carin “Multi-Label Learning from Medical Plain Text with Convolutional Residual Models”, *Machine Learning for Healthcare (MLHC)*, 2018 **Spotlight**

4. Y. Pu, S. Dai, **Z. Gan**, W. Wang, G. Wang, Y. Zhang, R. Henao and L. Carin “JointGAN: Multi-Domain Joint Distribution Learning with Generative Adversarial Nets”, *Int. Conf. Machine Learning (ICML)*, 2018
5. T. Xu, P. Zhang, Q. Huang, H. Zhang, **Z. Gan**, X. Huang and X. He “AttnGAN: Fine-Grained Text to Image Generation with Attentional Generative Adversarial Networks”, *Computer Vision and Pattern Recognition (CVPR)*, 2018
6. W. Wang, **Z. Gan**, W. Wang, D. Shen, J. Huang, W. Ping, S. Satheesh and L. Carin “Topic Compositional Neural Language Model”, *Artificial Intelligence and Statistics (AISTATS)*, 2018
7. Y. Pu, M. R. Min, **Z. Gan** and L. Carin “Adaptive Feature Abstraction for Translating Video to Text”, *Proc. American Association of Artificial Intelligence (AAAI)*, 2018

2017

1. **Z. Gan***, L. Chen*, W. Wang, Y. Pu, Y. Zhang, H. Liu, C. Li and L. Carin “Triangle Generative Adversarial Networks”, *Neural Information Processing Systems (NeurIPS)*, 2017
2. Y. Pu, W. Wang, R. Henao, L. Chen, **Z. Gan**, C. Li, and L. Carin “Adversarial Symmetric Variational Autoencoder”, *Neural Information Processing Systems (NeurIPS)*, 2017
3. Y. Pu, **Z. Gan**, R. Henao, C. Li, S. Han and L. Carin “VAE Learning via Stein Variational Gradient Descent”, *Neural Information Processing Systems (NeurIPS)*, 2017
4. Y. Zhang, D. Shen, G. Wang, **Z. Gan**, R. Henao and L. Carin “Deconvolutional Paragraph Representation Learning”, *Neural Information Processing Systems (NeurIPS)*, 2017
5. **Z. Gan**, Y. Pu, R. Henao, C. Li, X. He and L. Carin “Learning Generic Sentence Representations Using Convolutional Neural Networks”, *Conf. on Empirical Methods in Natural Language Processing (EMNLP)*, 2017 **Oral**
6. Y. Zhang, **Z. Gan**, K. Fan, Z. Chen, R. Henao, D. Shen and L. Carin “Adversarial Feature Matching for Text Generation”, *Int. Conf. Machine Learning (ICML)*, 2017
7. Y. Zhang, C. Chen, **Z. Gan**, R. Henao and L. Carin “Stochastic Gradient Monomial Gamma Sampler”, *Int. Conf. Machine Learning (ICML)*, 2017
8. **Z. Gan***, C. Li*, C. Chen, Y. Pu, Q. Su and L. Carin “Scalable Bayesian Learning of Recurrent Neural Networks for Language Modeling”, *Association for Computational Linguistics (ACL)*, 2017 **Oral**
9. **Z. Gan**, C. Gan, X. He, Y. Pu, K. Tran, J. Gao, L. Carin and L. Deng “Semantic Compositional Networks for Visual Captioning”, *Computer Vision and Pattern Recognition (CVPR)*, 2017 **Spotlight**
10. C. Gan, **Z. Gan**, X. He, J. Gao and L. Deng “StyleNet: Generating Attractive Visual Captions with Styles”, *Computer Vision and Pattern Recognition (CVPR)*, 2017
11. **Z. Gan**, P. D. Singh, A. Joshi, X. He, J. Chen, J. Gao and L. Deng “Character-level Deep Conflation for Business Data Analytics”, *Int. Conf. Acoustics, Speech and Signal Processing (ICASSP)*, 2017
12. Y. Xian, Y. Pu, **Z. Gan**, L. Lu and A. Thompson “Adaptive DCTNet for Audio Signal Classification”, *Int. Conf. Acoustics, Speech and Signal Processing (ICASSP)*, 2017
13. Q. Su, X. Liao, C. Li, **Z. Gan** and L. Carin “Unsupervised Learning with Truncated Gaussian Graphical Models”, *Proc. American Association of Artificial Intelligence (AAAI)*, 2017 **Oral**

2016

1. Y. Zhang, **Z. Gan** and L. Carin “Generating Text via Adversarial Training”, *NeurIPS Workshop*, 2016
2. Y. Xian, Y. Pu, **Z. Gan**, L. Lu and A. Thompson “Modified DCTNet for Audio Signals Classification”, *Journal of the Acoustical Society of America*, 2016
3. Y. Pu, **Z. Gan**, R. Henao, X. Yuan, C. Li, A. Stevens and L. Carin “Variational Autoencoder for Deep Learning of Images, Labels and Captions”, *Neural Information Processing Systems (NeurIPS)*, 2016
4. J. Song, **Z. Gan** and L. Carin “Factored Temporal Sigmoid Belief Networks for Sequence Learning”, *Int. Conf. Machine Learning (ICML)*, 2016

5. C. Li, A. Stevens, C. Chen, Y. Pu, **Z. Gan** and L. Carin "Learning Weight Uncertainty with Stochastic Gradient MCMC for Shape Classification", *Computer Vision and Pattern Recognition (CVPR)*, 2016 **Spotlight**
6. C. Chen, D. Carlson, **Z. Gan**, C. Li and L. Carin "Bridging the Gap Between Stochastic Gradient MCMC and Stochastic Optimization", *Artificial Intelligence and Statistics (AISTATS)*, 2016 **Oral**

2015

1. **Z. Gan**, C. Li, R. Henao, D. Carlson and L. Carin "Deep Temporal Sigmoid Belief Networks for Sequence Modeling", *Neural Information Processing Systems (NeurIPS)*, 2015
2. R. Henao, **Z. Gan**, J. Lu and L. Carin "Deep Poisson Factor Modeling", *Neural Information Processing Systems (NeurIPS)*, 2015
3. **Z. Gan**, C. Chen, R. Henao, D. Carlson and L. Carin "Scalable Deep Poisson Factor Analysis for Topic Modeling", *Int. Conf. Machine Learning (ICML)*, 2015
4. **Z. Gan**, R. Henao, D. Carlson and L. Carin "Learning Deep Sigmoid Belief Networks with Data Augmentation", *Artificial Intelligence and Statistics (AISTATS)*, 2015

Book Chapter

1. **Z. Gan**, X. Yuan, R. Henao, E. Tsalik and L. Carin "Inference of Gene Networks Associated with the Host Response to Infectious Disease", Chapter 13 of Book *Big Data Over Networks*. Cambridge University Press. In Press.

PhD Dissertation

1. **Z. Gan** "Deep Generative Models for Vision and Language Intelligence", Duke University.

Tutorial and Workshop

Tutorial

1. **Z. Gan**, L. Yu, Y. Cheng, L. Zhou, L. Li, Y.-C. Chen, J. Liu and X. He "Recent Advances in Vision-and-Language Research", *Computer Vision and Pattern Recognition (CVPR)*, 2020

Professional Activities

Area Chair: ICLR 2021, NeurIPS 2020, NeurIPS 2019

Senior Program Committee (SPC) Member: AAAI 2020

Interactive Arts Chair: ACMMM 2020

Conference Reviewer/PC Member:

- 2020: ICML, ICLR, IJCAI; CVPR, ECCV, ACMMM; ACL, EMNLP, COLING, AACL, CoNLL
- 2019: ICML, ICLR, AAAI; CVPR, ICCV, ACMMM; EMNLP, CoNLL
- 2018: NeurIPS, EMNLP, CVPR, ACCV
- 2016: NIPS

Journal Reviewer:

- Transactions on Pattern Analysis and Machine Intelligence (PAMI)
- Transactions on Image Processing (TIP)
- Transactions on Knowledge and Data Engineering (KDE)
- Journal of Selected Topics in Signal Processing (STSP)
- Transactions on Multimedia Computing Communications and Applications (TOMM)
- Transactions on Audio, Speech and Language Processing (ASL)
- Science China
- Transactions on Cybernetics, IET Computer Vision, Entropy

Workshop Reviewer/PC Member:

- 2020: ACL Workshop on Advances in Language and Vision Research
- 2019: ICCV Workshop on Closing the loop between Vision and Language
- 2019: ICLR Workshop on Deep Generative Models for Highly Structured Data
- 2018: ICML Workshop on Theoretical Foundations and Applications of Deep Generative Models

Talks

- “Deep Generative Models for Vision and Language Intelligence”, *Ph.D. Final Defense*, Durham, NC, February 2018
- “Deep Generative Models for Vision and Language Intelligence”, IBM Thomas J. Watson Research Center, Yorktown, NY, October 2017
- “Deep Generative Models for Vision and Language Intelligence”, NVIDIA, Santa Clara, CA, September 2017
- “Deep Generative Models for Vision and Language Intelligence”, Apple, Cupertino, CA, September 2017
- “Learning Generic Sentence Representations Using Convolutional Neural Networks”, *EMNLP*, Copenhagen, Denmark, September 2017
- “Semantic Compositional Networks for Visual Captioning”, *CVPR*, Hawaii, July 2017
- “Semantic Compositional Networks for Visual Captioning”, *Ph.D. Preliminary Exam*, Durham, NC, April 2017
- “Deep Generative Models for Sequence Learning”, *Ph.D. Qualifying Exam*, Durham, NC, December 2015

Competitions

- 2019/09: Rank 1st in GLUE benchmark
- 2019/06: Rank 2nd in Visual Dialog Challenge 2019
- 2018/09: Rank 3rd in Visual Dialog Challenge 2018

Software Skills

Python (Theano, Tensorflow, PyTorch), Matlab, R and C

Awards

AAAI-2020 Outstanding SPC Award.
ECE Fellowship, Duke University, 2013.
National Scholarship, Department of Minister of Education of China, 2010-2013.

Graduate Coursework

Bayesian and Modern Statistics, Probabilistic Machine Learning, Advanced Machine Learning, Statistical Inference, Statistical Computation, Information Theory, Graphical Models & Inference, Optimization For Engineers

Teaching Experience

Teaching Assistant	09/2014-12/2014
STA 601 - Bayesian and Modern Statistics	
Instructor: David Dunson, Ph.D	
Teaching Assistant	01/2015-05/2015
ECE 587 - Information Theory	
Instructor: Ahmad Beirami, Ph.D	